

DEEP LEARNING

WITH New mVITh

PHILIPP GABLER

2022-06-28

DEEP LEARNING

WITH ~~NO~~ MATH

AS LITTLE AS POSSIBLE

PHILIPP GABLER

2022-06-28




<https://xkcd.com/1838/>

OUTLINE

- ① A BIT OF HISTORY
- ② WHAT PROBLEM(S) DOES DL SOLVE?
/ COMPUTATIONAL
- ③ HOW DO SOLUTIONS LOOK LIKE?
/ ALGORITHMIC
- ④ HOW ARE SOLUTIONS IMPLEMENTED?
/ IMPLEMENTATIONAL
- ⑤ SOME LANGUAGE-RELATED USE CASES

Ev ăpxî q̂v tò PERCEPTRON



o MCCULLOCH & PITTS, 1948: NEURON MODEL



o ROSENBLATT, 1958: ENTHUSIASM

o MINSKY & PAPERT, 1969: SHOCK OF "PERCEPTRONS"

❄️ AI WINTER ❄️



PROBLEM OF LINEAR SEPARABILITY


NEW SPRING

- RUMELHART & MCCLELLAND (1986): PDP

- ↳ CONNECTIONISM

- ↳ BACKPROPAGATION

- ↳ SINGLE-LAYER → MULTI-LAYER



- STATISTICAL LEARNING / MACHINE LEARNING
OVER (SYMBOLIC) AI




TOWARDS DEEP LEARNING

- MORE LAYERS

- MORE DATA

- MORE COMPUTE (MOORE'S LAW)

- BETTER TUNING & TRAINING



TAKE-AWAYS 1


- ARTIFICIAL - NO NEURONS ANY MORE!
- AI AND ML HAVE SPLIT - DL \neq INTELLIGENCE
- GROWTH OF TECHNOLOGY IMPORTANT
- DL IS SUB SYMBOLIC (\approx CONNECTIONISM)

②

WHAT PROBLEMS DOES
DEEP LEARNING
SOLVE?


CHARACTERISTICS OF DEEP LEARNING

HIGH-DIMENSIONAL





PREDICTION


NON-LINEAR



ASSOCIATIVE




END-TO-END TRAINING,
PARAMETRIC



CLASSIFICATION



INPUT




MODEL



f



REGRESSION



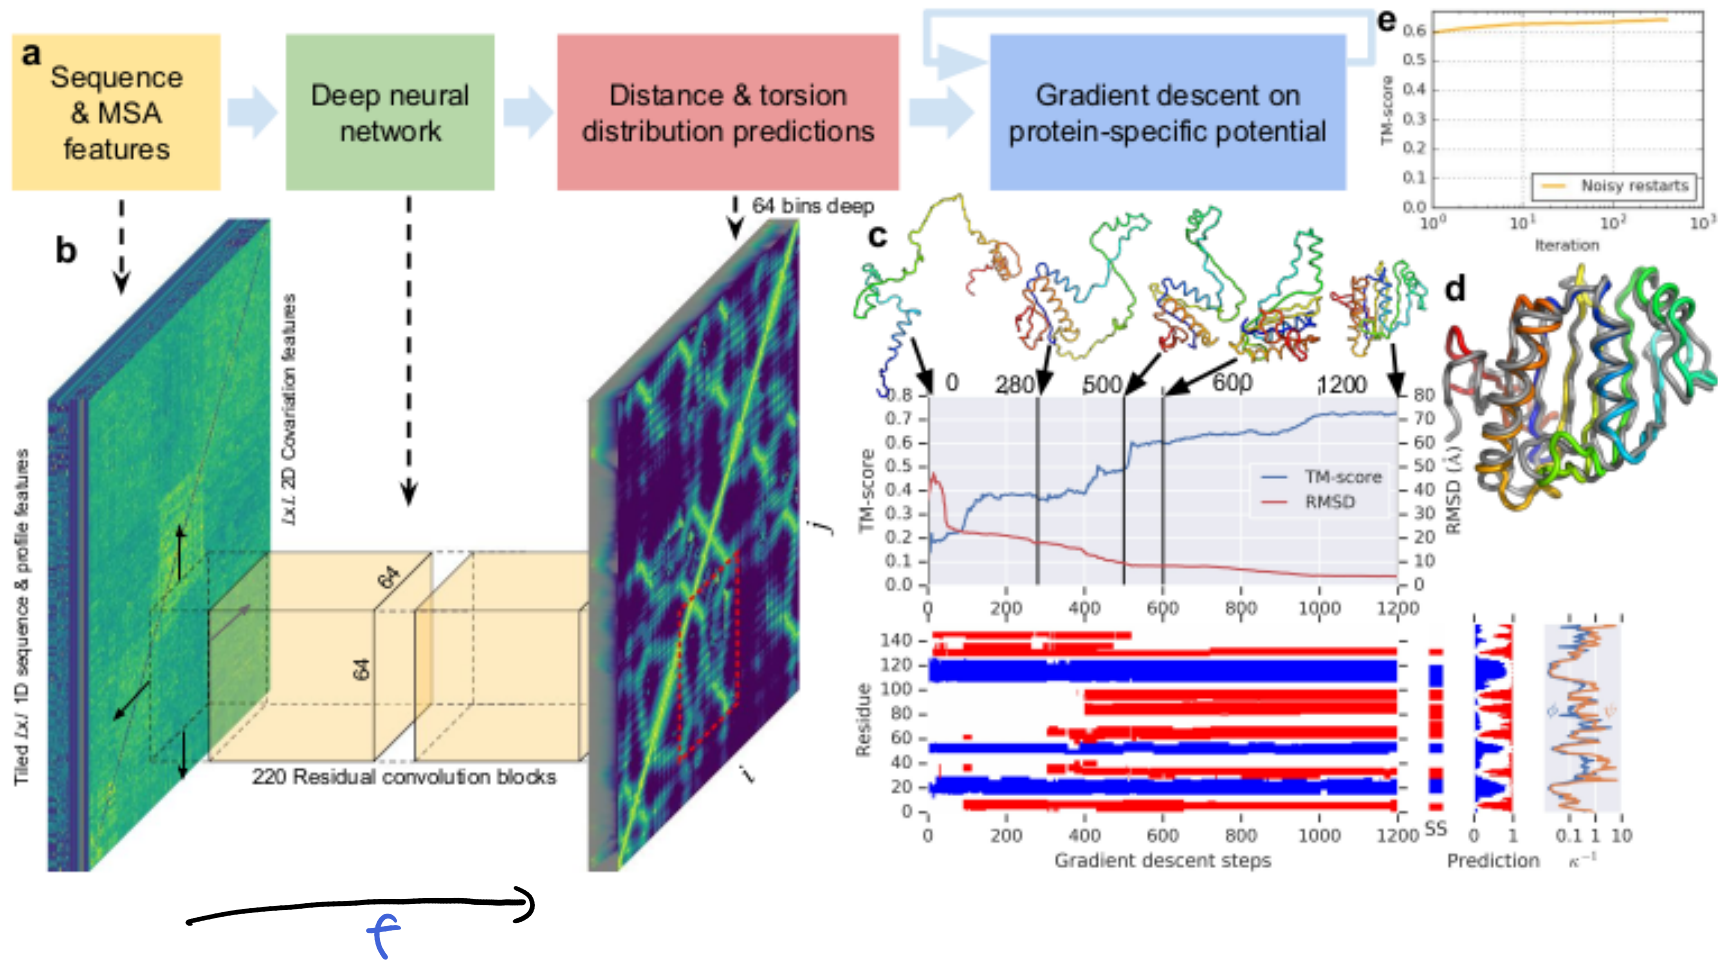
SEQUENCE PREDICTION

EI KAIKKI MUUMIT LAAKSOSSA



NICHT ALLE TASSEN IM SCHRANK

HIGH-DIMENSIONAL MODELLING




<https://doi.org/10.1038/s41586-019-1923-7>

TENSORS




JUST THINGS WITH MANY INDICES:



EVERYTHING IS LINEAR ALGEBRA!

TENSOR SHAPES

RANK 1




N-DIMENSIONAL
POINT



$[2.0, -0.3]$

RANK 2



B/W



$\begin{bmatrix} 0 & 0 & \dots & 0 \\ 0 & 1 & 1 & \dots & 0 \\ \vdots & & & & \vdots \\ 0 & \dots & & & 0 \end{bmatrix}$

x

RANK 3




RGB



$\begin{bmatrix} 0 & 1 & \dots & 0 \\ 0 & 2 & \dots & 0 \\ 0 & 0 & \dots & 0 \\ 0 & 1 & 1 & \dots & 0 \\ \vdots & & & & \vdots \\ 0 & \dots & & & 0 \end{bmatrix}$

R B

RANK 4



BATCH



$\begin{bmatrix} \dots \end{bmatrix}$

1

$\begin{bmatrix} \dots \end{bmatrix}$

2


$\begin{bmatrix} \dots \end{bmatrix}$

3

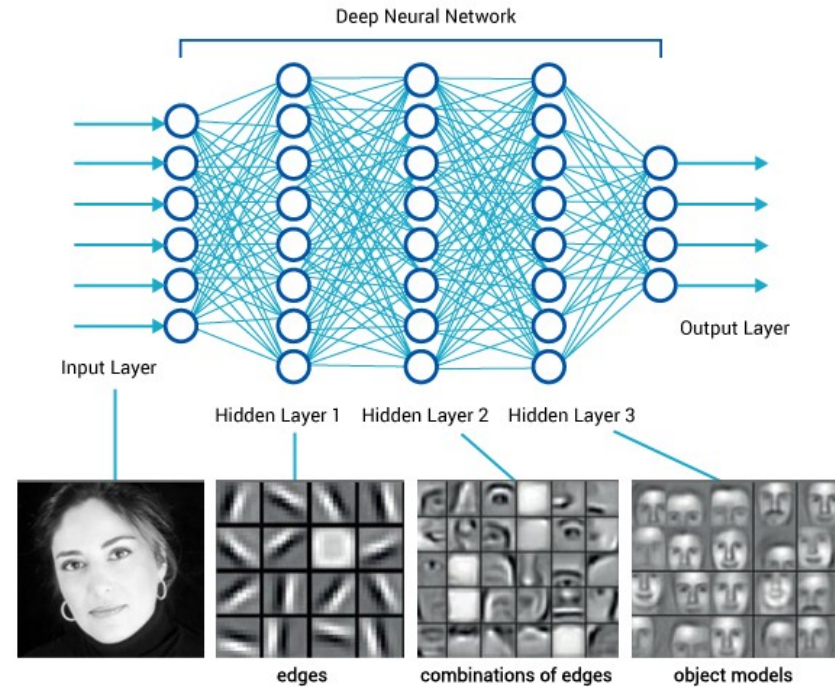
R B


FROM FEATURES TO LAYERS

TRADITIONAL






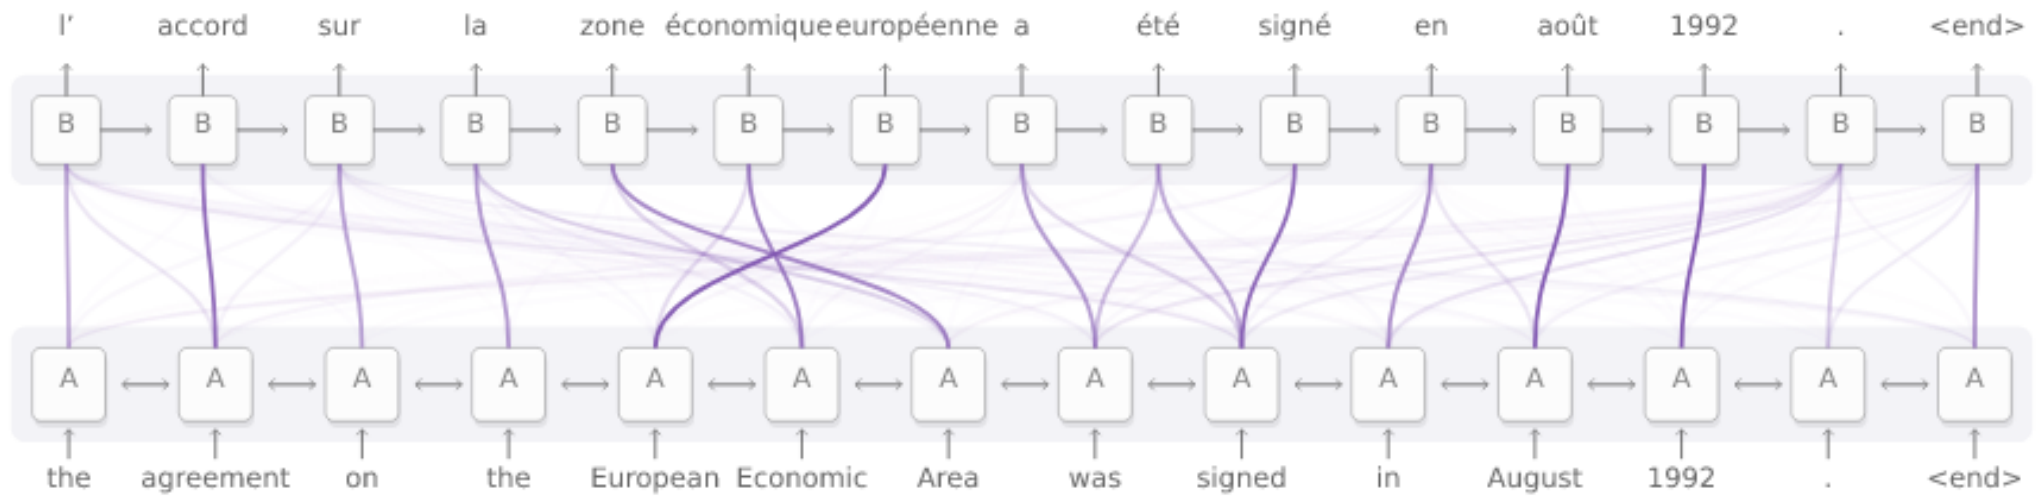
DEEP LEARNING






SPACE FOLDING





distill.pub/2016/augmented-rnns/



distill.pub/2016/augmented-rnns/

REFERENCES

INTRODUCTORY:

- Buckner, C. (2019). Deep learning: A philosophical introduction. Philosophy Compass, 14(10), e12625. <https://doi.org/10.1111/phc3.12625>
- Bzdok, D., Altman, N., & Krzywinski, M. (2018). Statistics versus machine learning. Nature Methods, 15(4), 233.
- Liquet, B., Moka, S., Nazarathy, Y. (2022). The Mathematical Engineering of Deep Learning (draft): Introduction. <https://deeplearningmath.org/drafts/chap1.pdf>

GENERAL:

- Bishop, C. M. (2006). Pattern recognition and machine learning. Springer.
- Murphy, K. P. (2012). Machine learning: A probabilistic perspective. MIT Press.
- Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.
<http://www.deeplearningbook.org>

REFERENCES

OPTIMIZATION & LEARNING:

- Kochenderfer, M. J., & Wheeler, T. A. (2019). Algorithms for Optimization. MIT Press.
- LeCun, Y., Chopra, S., Hadsell, R., Ranzato, M., & Huang, F. J. (2006). A Tutorial on Energy-Based Learning. <http://yann.lecun.org/exdb/publis/pdf/lecun-06.pdf>